
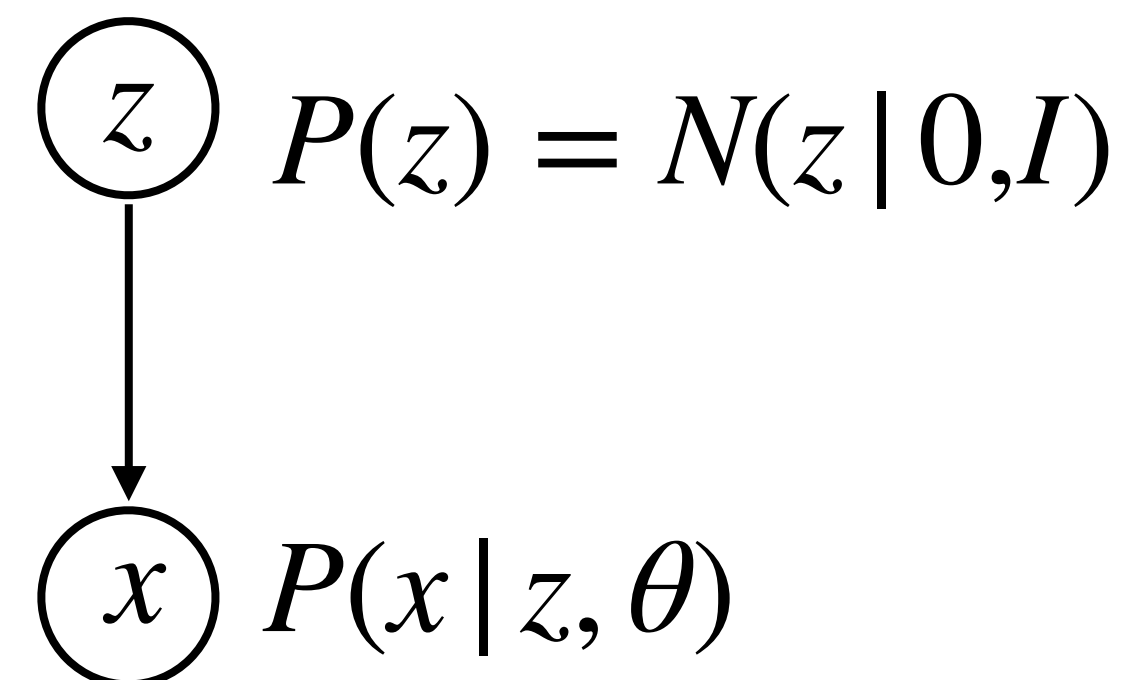


6.7900 Machine Learning (Fall 2024)

Lecture 22: generative models — diffusion

Preface: from simple samples to complex objects

- We can realize complex distributions over (e.g.) images by drawing samples from a fixed simple distribution and then mapping such samples through a complicated function (a neural network)



Preface: latents and how they are used

- The latent “degrees of freedom” z are useful if they are strongly coupled with x values; specifically, they should help us realize x ’s more easily conditionally on z
- In VAEs, we have an encoder and a decoder and both models represent this coupling between x and z .



- The ELBO criterion adjusts these models so they agree as distributions

Preface: latents and how they are used

- The latent “degrees of freedom” z are useful if they are strongly coupled with x values; specifically, they should help us realize x ’s more easily conditionally on z
- In VAEs, we have an encoder and a decoder and both models represent this coupling between x and z .




- The ELBO criterion adjusts these models so they agree as distributions, including marginals

$$\int Q_D(x) Q(z | x, \phi) dx \approx P(z) \quad \int P(z) P(x | z, \theta) dz \approx Q_D(x)$$


Preface: latents and how they are used

- We can actually always just **fix the encoder** and only learn the decoder



Preface: latents and how they are used

- We can actually always just **fix the encoder** and only learn the decoder




- The fixed encoder should still agree with the decoder's fixed latent marginal (otherwise they'd be permanently inconsistent)

$$\int Q_D(x)Q(z|x)dx \approx P(z)$$

Preface: latents and how they are used

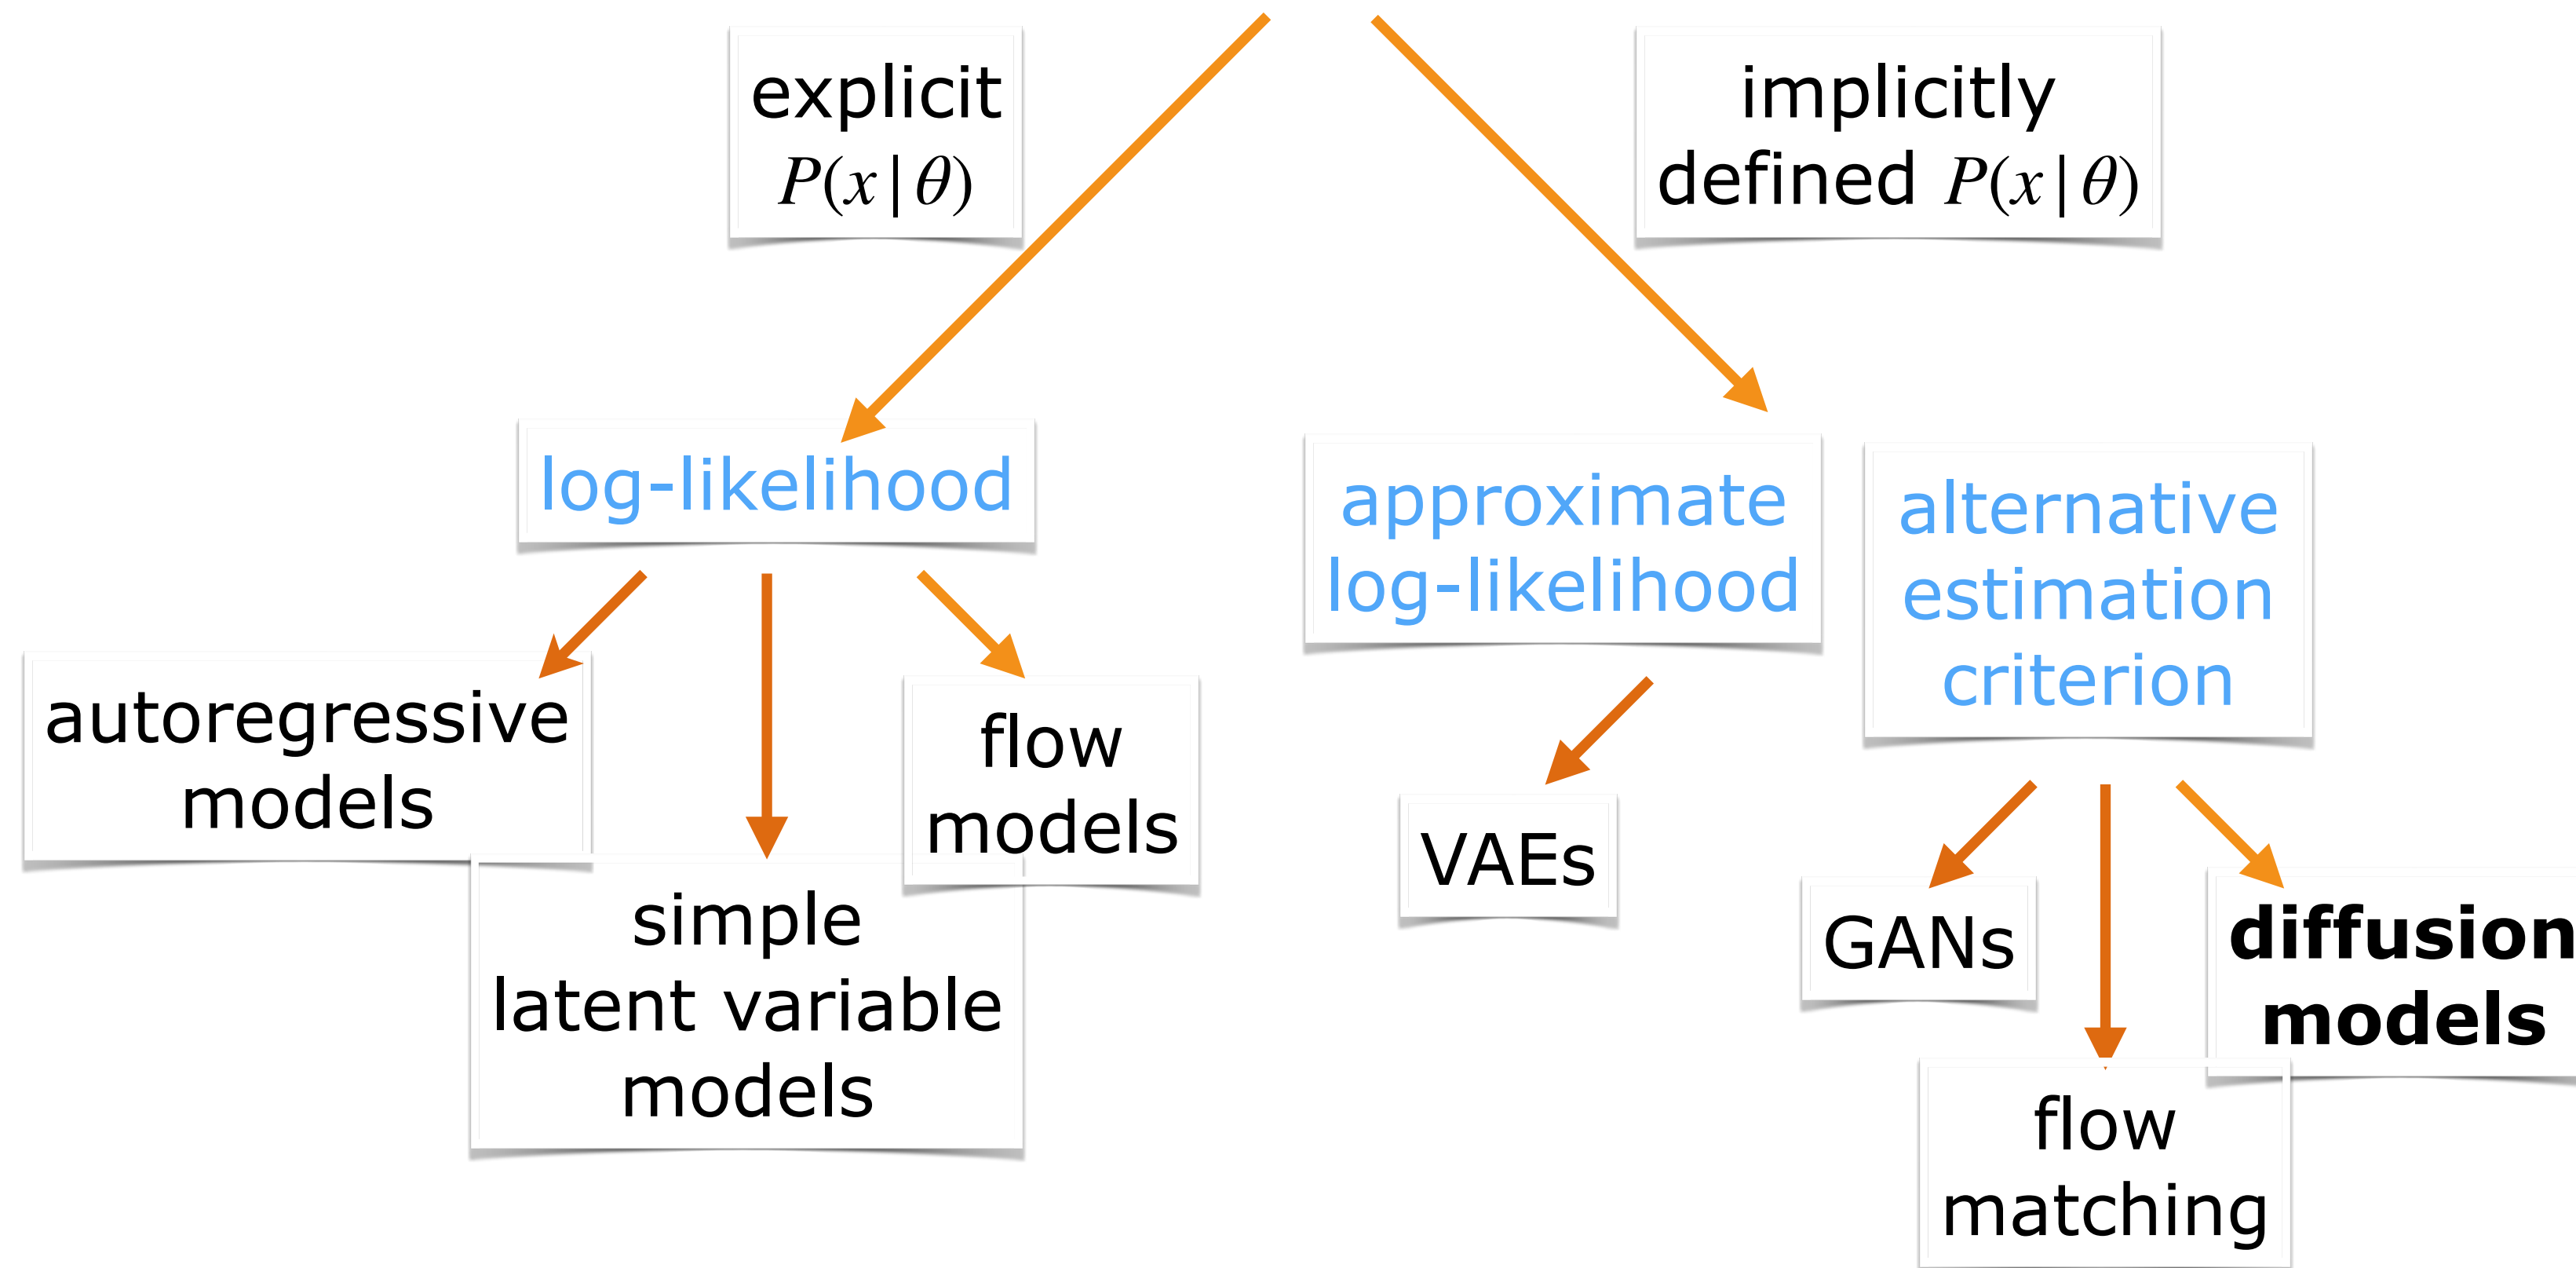
- We can actually always just **fix the encoder** and only learn the decoder



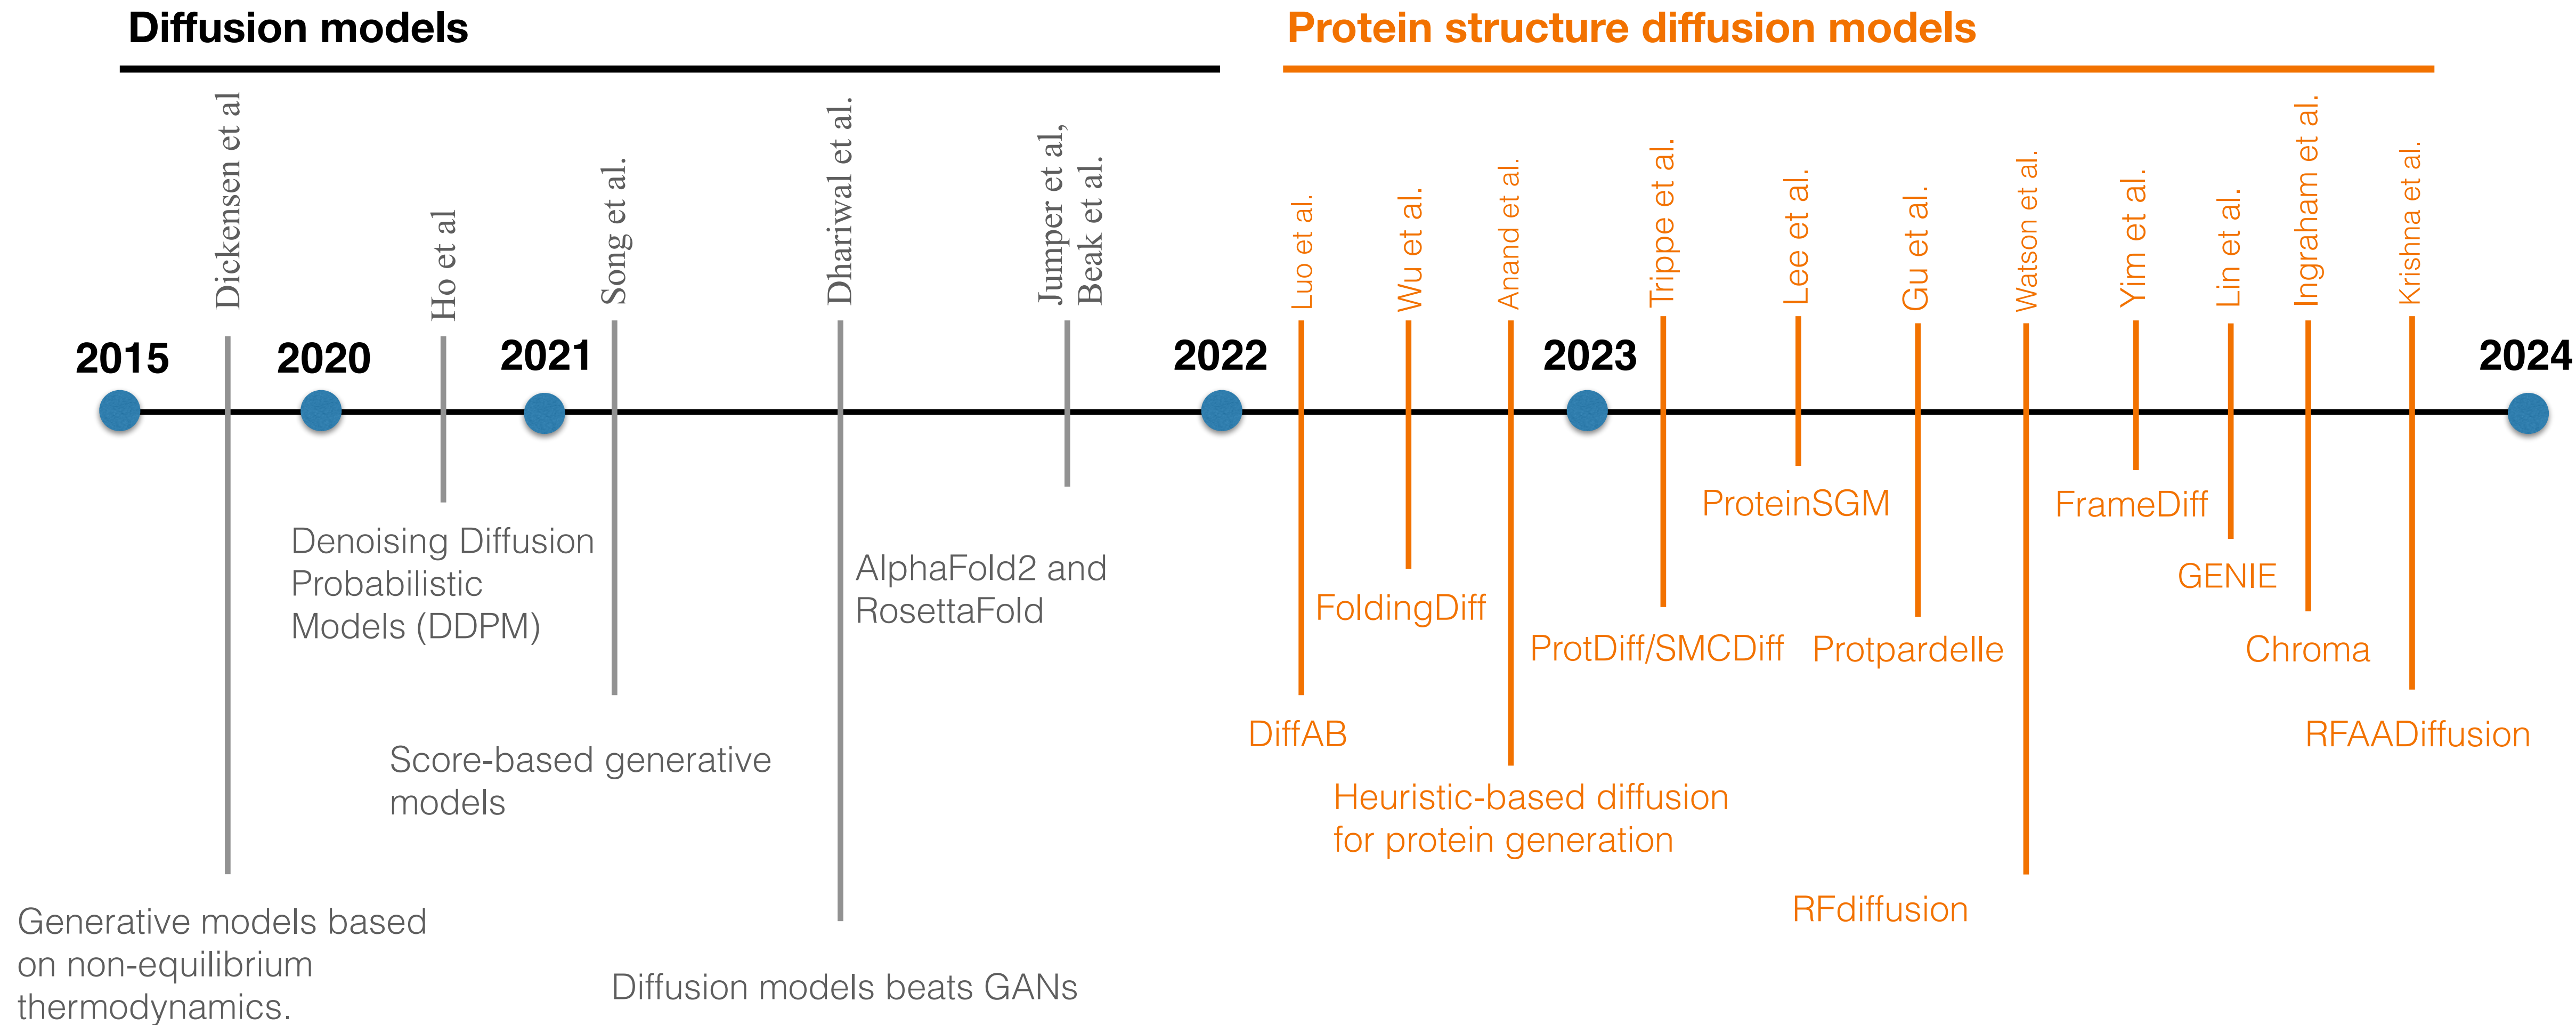
- The fixed encoder should still agree with the decoder's fixed latent marginal (otherwise they'd be permanently inconsistent)

$$\int Q_D(x)Q(z|x)dx \approx P(z)$$

A slice of the generative “landscape”




Ex: rapid adoption of diffusion models



[Yim et al 2024, "Diffusion models in protein structure and docking"]

Diffusion models


- De-noising diffusion models over images (e.g., Ho et al., Song et al.)



[image from
Rissanen et al 2022]


$$x \longleftarrow z \sim N(0, I)$$

- 3D structures of molecules (e.g., proteins)
- Discrete diffusion for language, sequence modeling, etc.



Diffusion motivation

- De-noising diffusion models over images (e.g., Ho et al., Song et al.)




[image from
Rissanen et al 2022]

x

- It's really helpful to give a generative model a stack of noisy versions of the same image... different features are present at different noise levels

Diffusion motivation

- De-noising diffusion models over images (e.g., Ho et al., Song et al.)



[image from
Rissanen et al 2022]

x


our model

- It's really helpful to give a generative model a stack of noisy versions of the same image... different features are present at different noise levels
- The goal is then to learn to successively uncover these lost features (de-noise images back into cleaner instances)

...Diffusion motivation



- De-noising diffusion models over images (e.g., Ho et al., Song et al.)



[image from
Rissanen et al 2022]



our model

- It's really helpful to give a generative model a stack of noisy versions of the same image... different features are present at different noise levels
- The goal is then to learn to successively uncover these features (de-noise images back into cleaner instances)
- This is hard since there are potentially many images that could have similar noisy versions

Diffusion motivation

- De-noising diffusion models over images (e.g., Ho et al., Song et al.)



[image from
Rissanen et al 2022]

$$x \longleftarrow z \sim N(0, I)$$

- It's really helpful to give a generative model a stack of noisy versions of the same image... different features are present at different noise levels
- The goal is then to learn to successively uncover these features (de-noise images back into cleaner instances)
- This is hard since there are potentially many images that could have similar noisy versions
- Once we have our de-noising model, we can draw a noisy sample, and iterate to generate a new clean instance (works the same with other types of objects, e.g., molecule structures)

Mechanisms of generation: diffusion


- De-noising diffusion models over images (e.g., Ho et al., Song et al.)



[image from
Rissanen et al 2022]

$$x_0 \longrightarrow z \sim N(0, I)$$

- A forward process “simplifies” objects (images) by adding more and more noise (each noise addition removes high frequency features from the image)



simple noise

Mechanisms of generation: diffusion

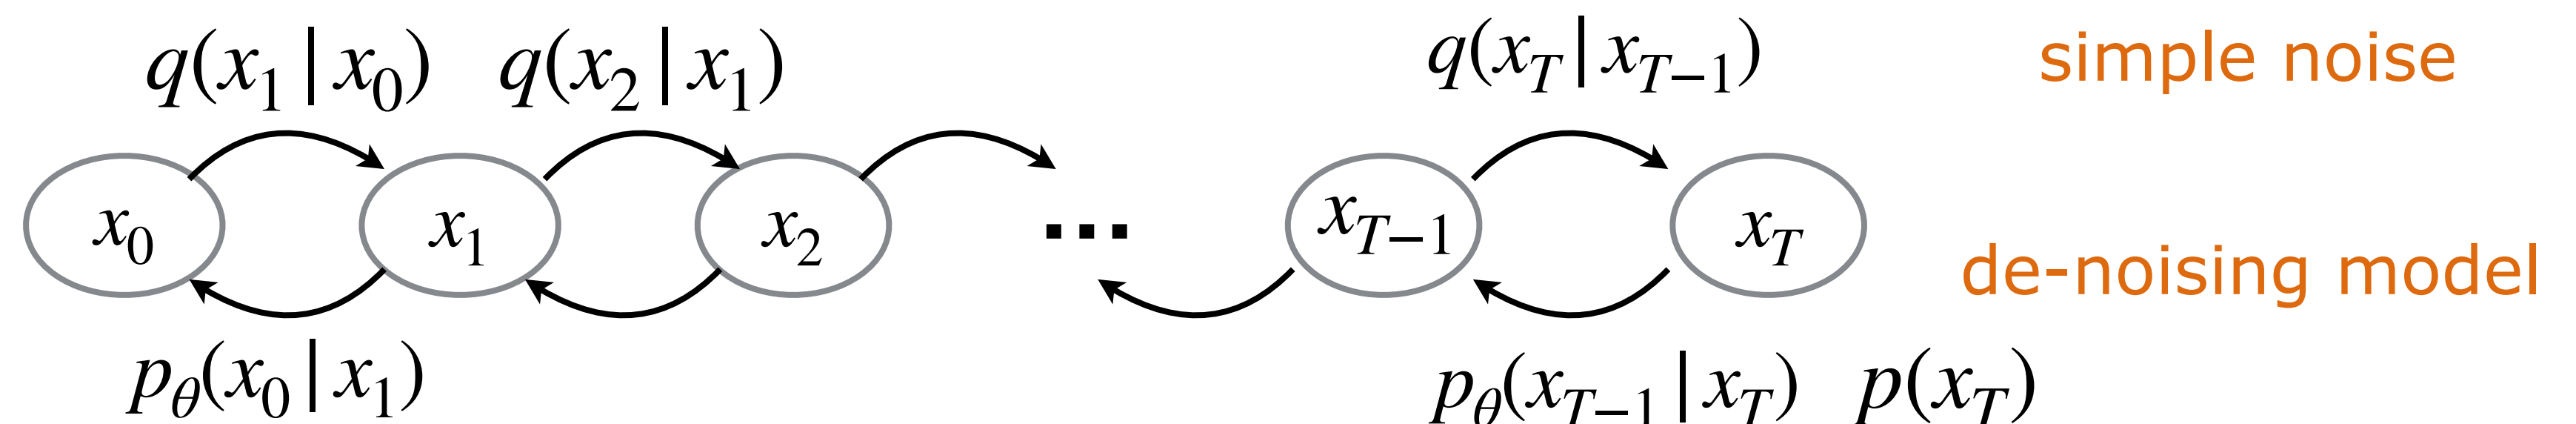
- De-noising diffusion models over images (e.g., Ho et al., Song et al.)




[image from
Rissanen et al 2022]

$$x_0 \longleftarrow z \sim N(0, I)$$

- A forward process “simplifies” objects (images) by adding more and more noise (each noise addition removes high frequency features from the image)
- The reverse process tries to reverse each step, i.e., it learns to predict how to de-noise images back into high quality versions




Basic diffusion models



$$q(x_{1:T}|x_0) = \begin{array}{c} q(x_1|x_0) \quad q(x_2|x_1) \quad \dots \quad q(x_T|x_{T-1}) \\ \begin{array}{c} \text{---} x_0 \text{---} x_1 \text{---} x_2 \text{---} \dots \text{---} x_{T-1} \text{---} x_T \text{---} \\ \text{---} p_\theta(x_0|x_1) \text{---} p_\theta(x_{T-1}|x_T) \text{---} p(x_T) \end{array} \end{array}$$

- We have a *fixed* (not learned) “forward process” $q(x_t|x_{t-1})$ which adds Gaussian noise (+ shrinkage)
- We use a reverse de-noising process $p_\theta(x_{t-1}|x_t)$ to (statistically) invert those steps
- This is a latent variable model where z = latent trajectories (x_1, \dots, x_T)

Basic diffusion models




$$q(x_{1:T}|x_0) = \begin{array}{c} q(x_1|x_0) \quad q(x_2|x_1) \quad \dots \quad q(x_T|x_{T-1}) \\ \begin{array}{c} \text{---} x_0 \text{---} x_1 \text{---} x_2 \text{---} \dots \text{---} x_{T-1} \text{---} x_T \text{---} \\ \text{---} p_\theta(x_0|x_1) \text{---} p_\theta(x_{T-1}|x_T) \text{---} p(x_T) \end{array} \end{array}$$

- We have a *fixed* (not learned) “forward process” $q(x_t|x_{t-1})$ which adds Gaussian noise (+ shrinkage)
- We use a reverse de-noising process $p_\theta(x_{t-1}|x_t)$ to (statistically) invert those steps
- This is a latent variable model where z = latent trajectories (x_1, \dots, x_T)
- In principle, we can learn the reverse de-noising process by maximizing the ELBO criterion

$$\log p_\theta(x_0) \geq E_q \left[\log \frac{p_\theta(x_{0:T})}{q(x_{1:T}|x_0)} \right]$$

Forward process




$$q(x_{1:T}|x_0) = q(x_1|x_0) q(x_2|x_1) \dots q(x_T|x_{T-1})$$

Diagram illustrating the forward process as a sequence of latent variables $x_0, x_1, x_2, \dots, x_{T-1}, x_T$ connected by arrows representing the conditional distributions $q(x_t|x_{t-1})$. The initial state x_0 is shaded gray, while subsequent states are white.

$$q(x_{1:T}|x_0) = \prod_{t=1}^T N(x_t | \sqrt{1 - \beta_t} x_{t-1}, \beta_t I)$$

- We start with x_0 and add a bit of Gaussian noise at each step (with shrinkage). So the distribution of x_t at step t will have to be Gaussian as well

Forward process: diffusion kernel



$$q(x_{1:T}|x_0) = q(x_1|x_0) q(x_2|x_1) \dots q(x_T|x_{T-1})$$

Diagram illustrating the forward process as a sequence of states $x_0, x_1, x_2, \dots, x_{T-1}, x_T$ connected by arrows representing the transition probabilities $q(x_t|x_{t-1})$. The initial state x_0 is shaded gray.


$$q(x_{1:T}|x_0) = \prod_{t=1}^T N(x_t | \sqrt{1 - \beta_t} x_{t-1}, \beta_t I)$$

- We start with x_0 and add a bit of Gaussian noise at each step (with shrinkage). So the distribution of x_t at step t will have to be Gaussian as well

$$q_t(x_t|x_0) = N(x_t | \sqrt{\bar{\alpha}_t} x_0, (1 - \bar{\alpha}_t)I) \quad (\text{diffusion kernel}) \quad \bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s)$$

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, \quad \epsilon \sim N(0, I)$$

Forward process: diffusion kernel



$$q(x_{1:T}|x_0) = q(x_1|x_0) q(x_2|x_1) \dots q(x_T|x_{T-1})$$

Diagram illustrating the forward process as a sequence of states $x_0, x_1, x_2, \dots, x_{T-1}, x_T$ connected by arrows representing the transition probabilities $q(x_t|x_{t-1})$. The initial state x_0 is shaded gray.


$$q(x_{1:T}|x_0) = \prod_{t=1}^T N(x_t | \sqrt{1 - \beta_t} x_{t-1}, \beta_t I)$$

- We start with x_0 and add a bit of Gaussian noise at each step (with shrinkage). So the distribution of x_t at step t will have to be Gaussian as well

$$q_t(x_t|x_0) = N(x_t | \sqrt{\bar{\alpha}_t} x_0, (1 - \bar{\alpha}_t)I) \quad (\text{diffusion kernel}) \quad \bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s)$$

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, \quad \epsilon \sim N(0, I) \quad \text{For large } T \quad x_T \approx \epsilon, \quad \epsilon \sim N(0, I)$$

Forward process: marginal at t




$$q(x_{1:T} | x_0) = \begin{array}{ccccccc} & q(x_1 | x_0) & q(x_2 | x_1) & & & q(x_T | x_{T-1}) & \\ & \curvearrowright & \curvearrowright & & & \curvearrowright & \\ (x_0) & (x_1) & (x_2) & \dots & & (x_{T-1}) & (x_T) \end{array}$$

- The diffusion kernel tells us how x_t is distributed conditioned on x_0

$$q_t(x_t | x_0) = N(x_t | \sqrt{\bar{\alpha}_t} x_0, (1 - \bar{\alpha}_t) I)$$

- If we sample data $x_0 \sim q(x_0)$ (e.g., uniform at random), the noisy image x_t at step t has a complex marginal distribution

Forward process: marginal at t




$$q(x_{1:T}|x_0) = q(x_1|x_0) q(x_2|x_1) \dots q(x_T|x_{T-1})$$

- The diffusion kernel tells us how x_t is distributed conditioned on x_0

$$q_t(x_t|x_0) = N(x_t | \sqrt{\bar{\alpha}_t}x_0, (1 - \bar{\alpha}_t)I)$$


- If we sample data $x_0 \sim q(x_0)$ (e.g., uniform at random), the noisy image x_t at step t has a complex marginal distribution

$$q_t(x_t) = \int q_t(x_t|x_0)q(x_0)dx_0 \approx \frac{1}{n} \sum_{i=1}^n q_t(x_t|x_0^i)$$



[Vahdat et al
2022]

Reverse process — our generative model




$$p_{\theta}(x_{0:T}) = \underbrace{p_{\theta}(x_0 | x_1)}_{\text{reverse process}} \underbrace{p_{\theta}(x_1 | x_2)}_{\text{reverse process}} \dots \underbrace{p_{\theta}(x_{T-1} | x_T)}_{\text{reverse process}} \underbrace{p(x_T)}_{\text{prior}}$$

- We use a reverse generative process to try to de-noise images back into clean versions, starting with x_T

$$p_{\theta}(x_{0:T}) = p(x_T) \prod_{t=1}^T N(x_{t-1} | \mu_{\theta}(x_t, t), \beta_t I)$$

Reverse process — our generative model



$$p_{\theta}(x_{0:T}) = p_{\theta}(x_0 | x_1) \quad \dots \quad p_{\theta}(x_{T-1} | x_T) \quad p(x_T)$$


The diagram shows a sequence of latent variables $x_0, x_1, x_2, \dots, x_{T-1}, x_T$ in ovals. Arrows point from x_1 to x_0 , from x_2 to x_1 , and from x_T to x_{T-1} . Ellipses between x_2 and x_{T-1} indicate intermediate steps in the sequence.

- We use a reverse generative process to try to de-noise images back into clean versions, starting with x_T

$$p_{\theta}(x_{0:T}) = p(x_T) \prod_{t=1}^T N(x_{t-1} | \mu_{\theta}(x_t, t), \beta_t I)$$

- We can set $p(x_T) = N(x_T | 0, I)$ since the forward process always ends with $x_T \sim N(0, I)$

Reverse process — our generative model



$$p_{\theta}(x_{0:T}) = p_{\theta}(x_0 | x_1) \quad \dots \quad p_{\theta}(x_{T-1} | x_T) \quad p(x_T)$$


The diagram shows a sequence of latent variables $x_0, x_1, x_2, \dots, x_{T-1}, x_T$ in circles. Arrows point from x_1 to x_0 , from x_2 to x_1 , and from x_T to x_{T-1} . Ellipses indicate the continuation of the sequence between x_2 and x_{T-1} .

- We use a reverse generative process to try to de-noise images back into clean versions, starting with x_T

$$p_{\theta}(x_{0:T}) = p(x_T) \prod_{t=1}^T N(x_{t-1} | \mu_{\theta}(x_t, t), \beta_t I)$$

- We can set $p(x_T) = N(x_T | 0, I)$ since the forward process always ends with $x_T \sim N(0, I)$
- Learning the de-noising “vector field” $\mu_{\theta}(x_t, t)$ is the key part!

Reparameterization of the reverse process



$$q(x_{1:T}|x_0) = \begin{array}{c} q(x_1|x_0) \quad q(x_2|x_1) \quad \dots \quad q(x_T|x_{T-1}) \\ \begin{array}{c} \text{Diagram showing a sequence of latent variables } x_0, x_1, x_2, \dots, x_{T-1}, x_T \text{ in ovals.} \\ \text{Forward arrows: } x_0 \rightarrow x_1 \rightarrow x_2 \rightarrow \dots \rightarrow x_{T-1} \rightarrow x_T \\ \text{Reverse arrows: } x_1 \rightarrow x_0, x_2 \rightarrow x_1, \dots, x_T \rightarrow x_{T-1} \end{array} \end{array}$$

$$p_\theta(x_{0:T}) = \begin{array}{c} p_\theta(x_0|x_1) \quad \dots \quad p_\theta(x_{T-1}|x_T) \quad p(x_T) \end{array}$$


$$p_\theta(x_{0:T}) = p(x_T) \prod_{t=1}^T N(x_{t-1} | \mu_\theta(x_t, t), \beta_t I)$$

$$\alpha_t = (1 - \beta_t)$$

$$\bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s)$$

Recall that based on the forward process: $x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, \quad \epsilon \sim N(0, I)$

Reparameterization of the reverse process



$$q(x_{1:T}|x_0) = \begin{array}{c} q(x_1|x_0) \quad q(x_2|x_1) \quad \dots \quad q(x_T|x_{T-1}) \\ \begin{array}{ccccccc} \textcircled{x_0} & \textcircled{x_1} & \textcircled{x_2} & \dots & \textcircled{x_{T-1}} & \textcircled{x_T} \\ \text{---} & \text{---} & \text{---} & \text{---} & \text{---} & \text{---} \\ \text{---} & \text{---} & \text{---} & \text{---} & \text{---} & \text{---} \end{array} \\ p_\theta(x_{0:T}) = \quad p_\theta(x_0|x_1) \quad \dots \quad p_\theta(x_{T-1}|x_T) \quad p(x_T) \end{array}$$


$$p_\theta(x_{0:T}) = p(x_T) \prod_{t=1}^T N(x_{t-1} | \mu_\theta(x_t, t), \beta_t I)$$

$$\alpha_t = (1 - \beta_t)$$

$$\bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s)$$

- Recall that based on the forward process: $x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon$, $\epsilon \sim N(0, I)$
- We can predict the added noise ϵ instead (at the same scale for all t).

Reparameterization of the reverse process



$$q(x_{1:T} | x_0) = \begin{array}{c} q(x_1 | x_0) \quad q(x_2 | x_1) \quad \dots \quad q(x_T | x_{T-1}) \\ \begin{array}{c} \text{Diagram showing a Markov chain } x_0 \rightarrow x_1 \rightarrow x_2 \rightarrow \dots \rightarrow x_{T-1} \rightarrow x_T \\ \text{with forward transitions } q(x_t | x_{t-1}) \text{ and reverse transitions } p_\theta(x_{t-1} | x_t) \end{array} \end{array}$$

$$p_\theta(x_{0:T}) = p_\theta(x_0 | x_1) \dots p_\theta(x_{T-1} | x_T) p(x_T)$$

$$p_\theta(x_{0:T}) = p(x_T) \prod_{t=1}^T N(x_{t-1} | \mu_\theta(x_t, t), \beta_t I)$$


$$\alpha_t = (1 - \beta_t)$$

$$\bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s)$$

- Recall that based on the forward process: $x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$, $\epsilon \sim N(0, I)$
- We can predict the added noise ϵ instead (at the same scale for all t). It allows us to calculate an estimate of x_0

$$\hat{x}_0 = \frac{1}{\sqrt{\bar{\alpha}}} (x_t - \sqrt{1 - \bar{\alpha}_t} \epsilon)$$

Reparameterization of the reverse process



$$q(x_{1:T} | x_0) = \begin{array}{c} q(x_1 | x_0) \quad q(x_2 | x_1) \quad \dots \quad q(x_T | x_{T-1}) \\ \begin{array}{c} \text{---} x_0 \text{---} x_1 \text{---} x_2 \text{---} \dots \text{---} x_{T-1} \text{---} x_T \text{---} \\ \text{---} p_\theta(x_0 | x_1) \text{---} p_\theta(x_{T-1} | x_T) \text{---} p(x_T) \end{array} \end{array}$$

$$p_\theta(x_{0:T}) = p(x_T) \prod_{t=1}^T \underline{N(x_{t-1} | \mu_\theta(x_t, t), \beta_t I)}$$

$$\alpha_t = (1 - \beta_t)$$


$$\bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s)$$

- Recall that based on the forward process: $x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$, $\epsilon \sim N(0, I)$
- We can predict the added noise ϵ instead (at the same scale for all t). It allows us to calculate an estimate of x_0 and what the de-noising step should be

$$\hat{x}_0 = \frac{1}{\sqrt{\bar{\alpha}_t}} (x_t - \sqrt{1 - \bar{\alpha}_t} \epsilon) \quad \underline{q(x_{t-1} | x_t, \hat{x}_0)} = \frac{q(x_t | x_{t-1}) q(x_{t-1} | \hat{x}_0)}{q(x_t | \hat{x}_0)}$$

all the
terms are
Gaussian

Reparameterization of the reverse process



$$q(x_{1:T}|x_0) = \begin{array}{c} q(x_1|x_0) \quad q(x_2|x_1) \quad \dots \quad q(x_T|x_{T-1}) \\ \begin{array}{c} \text{Diagram showing a Markov chain } x_0 \rightarrow x_1 \rightarrow x_2 \rightarrow \dots \rightarrow x_{T-1} \rightarrow x_T \text{ with forward arrows labeled } q(x_t|x_{t-1}) \text{ and backward arrows labeled } p_\theta(x_{t-1}|x_t). \end{array} \end{array}$$

$$p_\theta(x_{0:T}) = p_\theta(x_0|x_1) \dots p_\theta(x_{T-1}|x_T) p(x_T)$$

$$p_\theta(x_{0:T}) = p(x_T) \prod_{t=1}^T N(x_{t-1} | \mu_\theta(x_t, t), \beta_t I)$$


$$\alpha_t = (1 - \beta_t)$$

$$\bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s)$$

- Recall that based on the forward process: $x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$, $\epsilon \sim N(0, I)$
- We can predict the added noise ϵ instead (at the same scale for all t). It allows us to calculate an estimate of x_0 and what the de-noising step should be

$$\Rightarrow \mu_\theta(x_t, t) = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right)$$

Basic training, sampling




$$q(x_{1:T} | x_0) = \begin{array}{c} q(x_1 | x_0) \quad q(x_2 | x_1) \quad \dots \quad q(x_T | x_{T-1}) \\ \begin{array}{c} \text{---} x_0 \text{---} x_1 \text{---} x_2 \text{---} \dots \text{---} x_{T-1} \text{---} x_T \text{---} \\ \text{---} p_\theta(x_0 | x_1) \text{---} p_\theta(x_{T-1} | x_T) \text{---} p(x_T) \end{array} \end{array}$$

Algorithm 1 Training

- 1: **repeat**
 - 2: $\mathbf{x}_0 \sim q(\mathbf{x}_0)$
 - 3: $t \sim \text{Uniform}(\{1, \dots, T\})$
 - 4: $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
 - 5: Take gradient descent step on
$$\nabla_\theta \left\| \epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t) \right\|^2$$
 - 6: **until** converged
-

Basic training, sampling




$$q(x_{1:T} | x_0) = \begin{array}{c} q(x_1 | x_0) \quad q(x_2 | x_1) \quad \dots \quad q(x_T | x_{T-1}) \\ \begin{array}{c} \text{---} x_0 \text{---} x_1 \text{---} x_2 \text{---} \dots \text{---} x_{T-1} \text{---} x_T \text{---} \\ \text{---} p_\theta(x_0 | x_1) \text{---} p_\theta(x_{T-1} | x_T) \text{---} p(x_T) \end{array} \end{array}$$

strictly, ELBO would
imply a different
weighting

Algorithm 1 Training

- 1: **repeat**
 2. $\mathbf{x}_0 \sim q(\mathbf{x}_0)$
 - 3: $t \sim \text{Uniform}(\{1, \dots, T\})$
 - 4: $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
 - 5: Take gradient descent step on
$$\nabla_\theta \|\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t)\|^2$$
 - 6: **until** converged
-

Basic training, sampling



$$q(x_{1:T}|x_0) = \begin{matrix} & q(x_1|x_0) & q(x_2|x_1) & & & & q(x_T|x_{T-1}) \\ & \curvearrowright & \curvearrowright & & \curvearrowright & \curvearrowright & \\ x_0 & \xrightarrow{\quad} & x_1 & \xrightarrow{\quad} & x_2 & \cdots & x_{T-1} & \xrightarrow{\quad} & x_T \\ & \curvearrowleft & \curvearrowleft & & \curvearrowleft & \curvearrowleft & \\ & p_\theta(x_0|x_1) & & & & & p_\theta(x_{T-1}|x_T) & & p(x_T) \end{matrix}$$

strictly, ELBO would
imply a different
weighting

Algorithm 1 Training


- 1: **repeat**
- 2: $\mathbf{x}_0 \sim q(\mathbf{x}_0)$
- 3: $t \sim \text{Uniform}(\{1, \dots, T\})$
- 4: $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 5: Take gradient descent step on

$$\nabla_\theta \|\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t)\|^2$$
- 6: **until** converged

Algorithm 2 Sampling

- 1: $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 2: **for** $t = T, \dots, 1$ **do**
- 3: $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ if $t > 1$, else $\mathbf{z} = \mathbf{0}$
- 4: $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$
- 5: **end for**
- 6: **return** \mathbf{x}_0

Score based (continuous) diffusion models



x_0


x_1

x_{T-1}

x_T

- What if we used more steps in between $[0, T]$, adding less noise at each?


Score based (continuous) diffusion models



- What if we used more steps in between $[0, T]$, adding less noise at each? We can write the noise variance now as $\beta(t)\Delta t$ (approaching zero)

$$x_t = \sqrt{1 - \beta(t)\Delta t} x_{t-1} + \sqrt{\beta(t)\Delta t} \epsilon, \quad \epsilon \sim N(0, I)$$

Score based (continuous) diffusion models



x_0

x_1

x_{T-1}


x_T

- What if we used more steps in between $[0, T]$, adding less noise at each? We can write the noise variance now as $\beta(t)\Delta t$ (approaching zero)

$$x_t = \sqrt{1 - \beta(t)\Delta t} x_{t-1} + \sqrt{\beta(t)\Delta t} \epsilon, \quad \epsilon \sim N(0, I)$$

$$x_t - x_{t-1} = (\sqrt{1 - \beta(t)\Delta t} - 1) x_{t-1} + \sqrt{\beta(t)\Delta t} \epsilon, \quad \epsilon \sim N(0, I)$$

Score based (continuous) diffusion models



x_0

x_1

x_{T-1}

x_T


- What if we used more steps in between $[0, T]$, adding less noise at each? We can write the noise variance now as $\beta(t)\Delta t$ (approaching zero)

$$x_t = \sqrt{1 - \beta(t)\Delta t} x_{t-1} + \sqrt{\beta(t)\Delta t} \epsilon, \quad \epsilon \sim N(0, I)$$

$$x_t - x_{t-1} = (\sqrt{1 - \beta(t)\Delta t} - 1) x_{t-1} + \sqrt{\beta(t)\Delta t} \epsilon, \quad \epsilon \sim N(0, I)$$

$$\Delta x_t \approx -\frac{1}{2}\beta(t)\Delta t x_{t-1} + \sqrt{\beta(t)\Delta t} \epsilon, \quad \epsilon \sim N(0, I)$$

Score based (continuous) diffusion models



x_0

x_1

x_{T-1}

x_T

- What if we used more steps in between $[0, T]$, adding less noise at each? We can write the noise variance now as $\beta(t)\Delta t$ (approaching zero)

$$x_t = \sqrt{1 - \beta(t)\Delta t} x_{t-1} + \sqrt{\beta(t)\Delta t} \epsilon, \quad \epsilon \sim N(0, I)$$

$$x_t - x_{t-1} = (\sqrt{1 - \beta(t)\Delta t} - 1) x_{t-1} + \sqrt{\beta(t)\Delta t} \epsilon, \quad \epsilon \sim N(0, I)$$

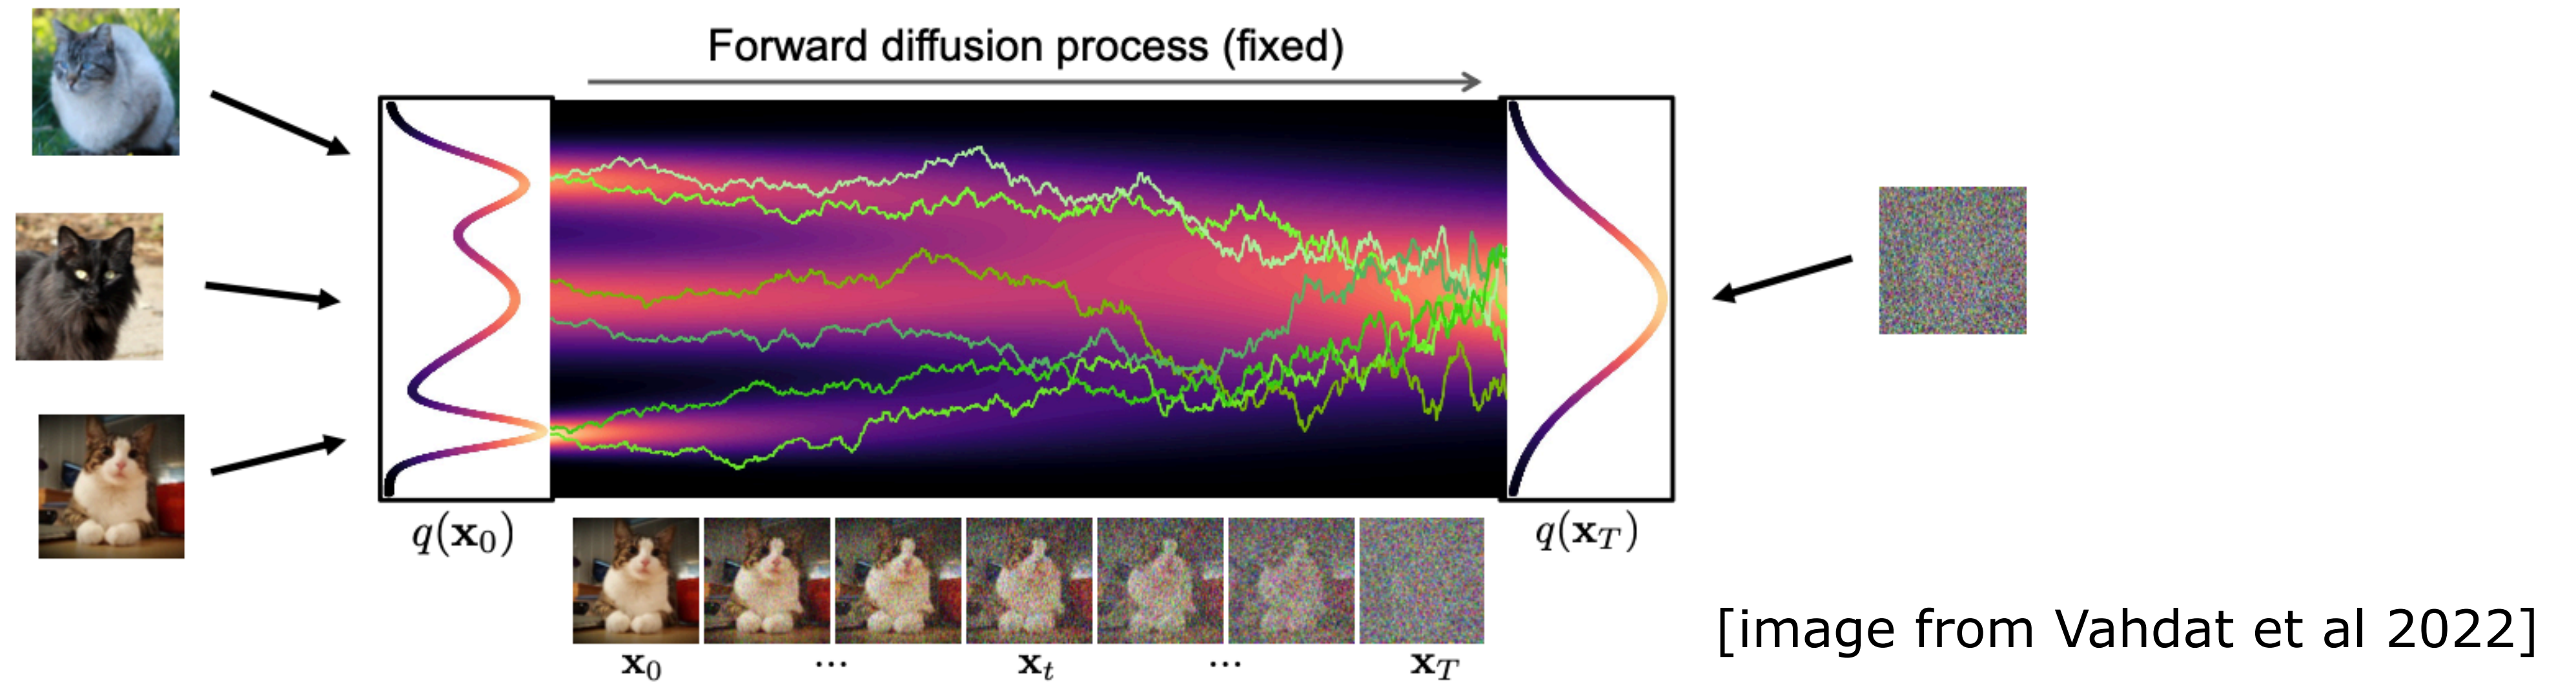
$$\Delta x_t \approx -\frac{1}{2}\beta(t)\Delta t x_{t-1} + \sqrt{\beta(t)\Delta t} \epsilon, \quad \epsilon \sim N(0, I)$$

- Taking the limit of Δt we arrive at a stochastic differential equation (SDE)

$$dx_t = -\frac{1}{2}\beta(t) x_t dt + \sqrt{\beta(t)} dw_t$$

Brownian motion

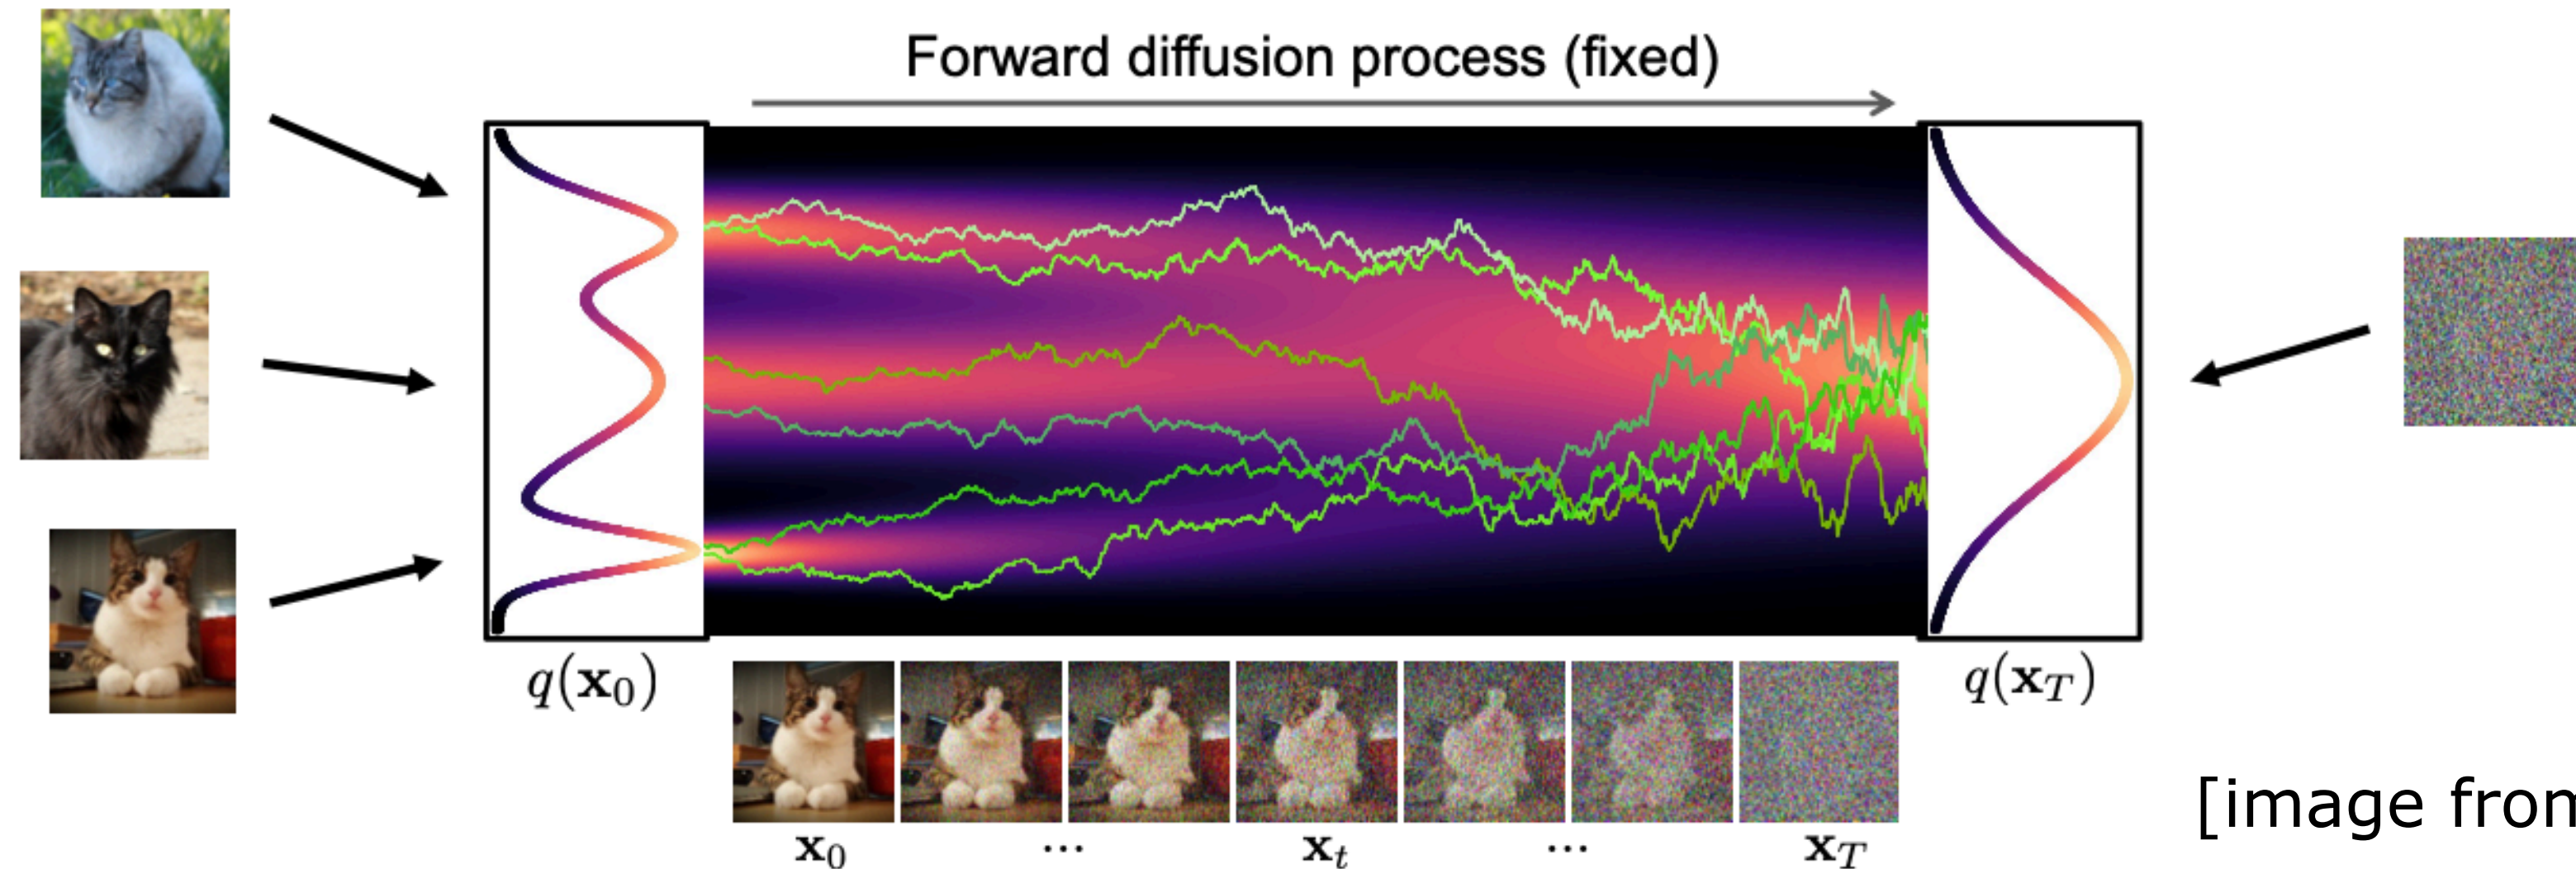
Forward diffusion process



Forward process:

$$dx_t = -\frac{1}{2}\beta(t)x_t dt + \sqrt{\beta(t)}dw_t$$

Reverse diffusion process



- Forward process: $dx_t = -\frac{1}{2}\beta(t)x_t dt + \sqrt{\beta(t)} dw_t$
- Key result [Anderson 82]: reverse process is also a diffusion process

$$dx_t = \left[-\frac{1}{2}\beta(t)x_t - \underbrace{\beta(t) \nabla_{x_t} \log q_t(x_t)}_{\text{score function}} \right] dt + \sqrt{\beta(t)} d\tilde{w}_t$$

also Brownian motion

(dt now negative)

Reverse diffusion process

- We would like to use the reverse process to sample new images $(x_T \sim N(0, I)$ as before)

$$dx_t = \left[-\frac{1}{2}\beta(t)x_t - \underbrace{\beta(t) \nabla_{x_t} \log q_t(x_t)}_{\text{score function}} \right] dt + \sqrt{\beta(t)} d\tilde{w}_t$$

- To do so we should learn a neural model $s_\theta(x, t)$ to approx. the score function
- How to learn $s_\theta(x, t)$?

Reverse diffusion process

- We would like to use the reverse process to sample new images $(x_T \sim N(0, I)$ as before)

$$dx_t = \left[-\frac{1}{2}\beta(t)x_t - \beta(t) \underbrace{\nabla_{x_t} \log q_t(x_t)}_{\text{score function}} \right] dt + \sqrt{\beta(t)} d\tilde{w}_t$$

- To do so we should learn a neural model $s_\theta(x, t)$ to approx. the score function
- How to learn $s_\theta(x, t)$? We could try to (similar to before)

$$x_0 \sim q(x_0), \quad t \sim U(0, T)$$

$$x_t \sim q_t(x_t | x_0) \quad \text{diffusion kernel (it is still just a Gaussian)}$$

$$\theta \leftarrow \theta - \eta \nabla_\theta \|s_\theta(x_t, t) - \underbrace{\nabla_{x_t} \log q_t(x_t)}\|^2$$

- But $q_t(x_t)$ is a complex distribution!

$$q_t(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t; \gamma_t \mathbf{x}_0, \sigma_t^2 \mathbf{I})$$

$$\gamma_t = e^{-\frac{1}{2} \int_0^t \beta(s) ds}$$

$$\sigma_t^2 = 1 - e^{-\int_0^t \beta(s) ds}$$

[Vahdat et al 2022]

Reverse diffusion process

- We would like to use the reverse process to sample new images $(x_T \sim N(0, I)$ as before)

$$dx_t = \left[-\frac{1}{2}\beta(t)x_t - \beta(t) \underbrace{\nabla_{x_t} \log q_t(x_t)}_{\text{score function}} \right] dt + \sqrt{\beta(t)} d\tilde{w}_t$$

- To do so we should learn a neural model $s_\theta(x, t)$ to approx. the score function
- How to learn $s_\theta(x, t)$? We can do

$$x_0 \sim q(x_0), \quad t \sim U(0, T)$$

$$x_t \sim q_t(x_t | x_0) \quad \text{diffusion kernel (it is still just a Gaussian)}$$

$$\theta \leftarrow \theta - \eta \nabla_\theta \|s_\theta(x_t, t) - \underbrace{\nabla_{x_t} \log q_t(x_t | x_0)}\|^2$$

$$q_t(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t; \gamma_t \mathbf{x}_0, \sigma_t^2 \mathbf{I})$$

$$\gamma_t = e^{-\frac{1}{2} \int_0^t \beta(s) ds}$$

$$\sigma_t^2 = 1 - e^{-\int_0^t \beta(s) ds}$$

[Vahdat et al 2022]

- $\nabla_{x_t} \log q_t(x_t | x_0)$ (score of a Gaussian) can be easily calculated and has the same expected value (over x_0 for a given x_t) when $x_0 \sim q(x_0)$!

Score expectation: derivation

- $\nabla_{x_t} \log q_t(x_t | x_0)$ is a score of a Gaussian that can be easily calculated
- It has the same expected value as the complex score function $\nabla_{x_t} \log q_t(x_t)$ we want and thus can be used as a (noisy) learning target

$$\begin{aligned} E_{x_0 \sim q(x_0 | x_t)} \{ \nabla_{x_t} \log q_t(x_t | x_0) \} &= \int q(x_0 | x_t) \nabla_{x_t} \log q_t(x_t | x_0) dx_0 \\ &= \int \frac{q(x_0) q_t(x_t | x_0)}{q_t(x_t)} \nabla_{x_t} \log q_t(x_t | x_0) dx_0 \\ &= \int \frac{q(x_0)}{q_t(x_t)} \nabla_{x_t} q_t(x_t | x_0) dx_0 \\ &= \frac{1}{q_t(x_t)} \nabla_{x_t} \int q_t(x_t | x_0) q(x_0) dx_0 \\ &= \frac{1}{q_t(x_t)} \nabla_{x_t} q_t(x_t) = \nabla_{x_t} \log q_t(x_t) \end{aligned}$$

Reverse diffusion process

- We would like to use the reverse process to sample new images $(x_T \sim N(0, I)$ as before)

$$dx_t = \left[-\frac{1}{2}\beta(t)x_t - \beta(t) \underbrace{\nabla_{x_t} \log q_t(x_t)}_{\text{score function}} \right] dt + \sqrt{\beta(t)} d\tilde{w}_t$$

- To do so we should learn a neural model $s_\theta(x, t)$ to approx. the score function
- How to learn $s_\theta(x, t)$? We could try to (similar to before)

$$x_0 \sim q(x_0), \quad t \sim U(0, T)$$

$$x_t \sim q_t(x_t | x_0) \quad \text{diffusion kernel (it is still just a Gaussian)}$$

$$\theta \leftarrow \theta - \eta \nabla_\theta \|s_\theta(x_t, t) - \underbrace{\nabla_{x_t} \log q_t(x_t)}\|^2$$

$$q_t(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t; \gamma_t \mathbf{x}_0, \sigma_t^2 \mathbf{I})$$

$$\gamma_t = e^{-\frac{1}{2} \int_0^t \beta(s) ds}$$

$$\sigma_t^2 = 1 - e^{-\int_0^t \beta(s) ds}$$

[Vahdat et al 2022]

Additional reading

- **Some early diffusion model papers:**
- Sohl-Dickstein et al., “Deep Unsupervised Learning using Nonequilibrium Thermodynamics”, <http://proceedings.mlr.press/v37/sohl-dickstein15.pdf>
- Ho et al., “Denoising Diffusion Probabilistic Models”, <https://arxiv.org/abs/2006.11239>
- Song et al., “Generative Modeling by Estimating Gradients of the Data Distribution”, <https://arxiv.org/abs/1907.05600>
- **Tutorials (excerpts used in this lecture)**
- A. Vahdat, K. Kreis, R. Gao, “CVPR 2022 Tutorial — Denoising Diffusion-based Generative Modeling: Foundations and Applications”, <https://cvpr2022-tutorial-diffusion-models.github.io/>